

# COSC 2P95

## Error checking and debugging

Week 10

Brock University

# Reliability

We've spent several weeks now learning to write solutions, using various different tools and approaches.

However, outside of minimal *sanity checking*, we've had to mostly rely on inputs/datas/solutions that “just work”.

- That really isn't plausible for a complete system of any significant size

There are a few basic concerns when developing solutions:

- Code must be written ‘correct’ to the language
- Algorithms should be free of logic errors
- Never trust a human
  - ▶ Never trust a machine, either, if you can avoid it
- When things go wrong (and they will):
  - ▶ Recovering is nice
  - ▶ *Detecting* the problem is paramount

# Syntax vs semantics

Before we get any further, we first need to ensure that we remember the difference between *syntax* and *semantics* — and, by extension, syntax errors vs semantic errors

- Forgetting a semicolon? That's a ~~padding~~ syntax error
- Using the wrong index, or calculating a formula wrong? Semantics!
- Generally, this is the difference between compile-time and runtime
  - ▶ Rule of thumb: compile errors are no big deal. Finding out something's broken after you've deployed it? Not pleasant

# What actually goes wrong?

## What do we do about it?

Oftentimes, outside of pure logic/algorithm errors, semantic errors are a result of *violated assumptions*.

- e.g. you tried reading from a stream; since it returned, you assume it didn't use a default 0
- Or you asked the user for a number from 1 to 10, and *assumed* they wouldn't enter "cheeseburger"

What can we do about it? Well, the first step is *defensive programming*.

- Don't write loops that end on zero, if you could end on  $\leq 0$
- For prompts for user input, consider putting them in loops until they receive legal input, or include a 'default behaviour'
- Check status flags of streams after access

However, we can still find ourselves making assumptions. Can we verify them?

# Assertions

A common debugging technique is to add dozens of *print* statements, to report *happiness messages*.

- If the code got to a certain point, it prints
- If a value contains what you expected, it prints

But, for any significant quantity of tests, that quickly becomes impossible to sift through. There's an alternative: *assertions*. e.g.:

```
#include <cassert>
...
int value;
std::cout<<"Give a number from 1 to 5: ";
std::cin>>value;
assert(value>=1 && value<=5);
std::cout<<"Good choice!"<<std::endl;
```

Let's try running this to see what happens.

## Better assertions

Sometimes, you might want to either define a `bool` for your test, or use the `assert` to check for *non-null*. For these cases, even if it failed, you'd end up with an output equivalent to 'infile' failed.

- If desired, you can write something like:  

```
assert(infile && "Much more descriptive feedback");
```

  - ▶ The status of `infile` is what will dictate the result

# Great! I'm done!

Should I keep the testing code in?

This may sound counterintuitive, but you *don't* want the assertions in your final release.

- The typical end-user won't be able to do anything with the error messages anyway
- The errors might not be catastrophic, but the program *won't* execute past tripped assertions
- Even with the previous slide's suggestion, the feedback typically isn't adequate for proper troubleshooting
- Though bugs are still possible, by the time you get to production, you should be confident that you've removed the major bugs that would have triggered your assertions
  - ▶ Things like “I hope I understand how this formula works” and “my file-loading function actually loads files” *should* be pretty reliable

If you add the preprocessor directive `#define NDEBUG` *before* including `cassert`, you turn them off.

- It's also slightly more efficient

# So, now my code is done?

I can trust its operation?

- *Ah ha ha ha ha ha ha*



... really?

*...ha ha ha ha ha ha ha!*

# Stoppit

Sorry.

No, we still have several concerns:

- Even completely correct code will sometimes encounter bad situations
  - ▶ Maybe we can recover; maybe not. But we need to identify them
  - ▶ We may need to handle such issues at a different level from the local context where they're encountered
- Some issues (entering a number outside the required range) can be immediately addressed, but how do we tell a calling context that there was an issue inside a function?
  - ▶ Return a `bool` type? Yes, sometimes
  - ▶ Include a reference parameter for a status flag? Possibly, but that gets nasty for multiple possible errors
  - ▶ For classes, add a status flag? Sure, if you're already using a class, and changing the specification is okay

Basically, we have a *lot* of other cases. What we really need is the ability to flag an issue, suspend execution, and immediately jump straight to code intended to handle that issue.

# Exceptional events

C++ offers a tool to us: the *exception*.

- We're *not* going to learn *goto* statements in this course, but as a thought experiment, suppose your program could respond to a problem by *immediately* dropping everything, and *jumping* straight to a special block of code, dedicated to handling that problem
- If that problem's never encountered, that jump never occurs, and that special code is never used

We *do* have a mechanism for this.

# Exceptions

## throw

In C++, we can throw whatever we like. An `int`, a `float`, a `Monkey` (but ignore that one for now).

- When something is *thrown*, execution at the current position in the code *immediately stops*
- We'll get into more detail in just a moment, but basically the program tries to find a *handler* capable of receiving whatever was thrown
- Depending on the situation, either the fact that *anything* was thrown, or precisely *what* was thrown will indicate the problem, and the program will respond appropriately
  - ▶ Responding could mean correcting the issue, or asking the user for additional input, or it could simply trigger a *graceful exit* of the program (closing streams, etc., and possibly displaying a message or error code)

# Exceptions

## try/catch

Simply throwing something wouldn't have any use. Instead, we'll *normally* want to tell the system to get ready for it.

- We put code that we anticipate potentially triggering an *exceptional event* inside a try block
- We follow a try block with a catch block
  - ▶ The catch block is matched up to the same type as what's thrown, and receives the thrown value
  - ▶ The catch block is where you handle that exceptional event
- A single try block may have *multiple* catch blocks
  - ▶ Each catch block is matched to a different thrown type
    - ★ If you want to throw, for example, multiple int types, you'll need a condition in your catch block

Can you throw a value *without* a try/catch? Yes, but it isn't advisable. Example time!

# Exceptions

## Different contexts/call frames

Of course, it would be ridiculous if exceptions needed to be thrown within the same context as their try/catch blocks.

- How would that be significantly different from any other regular conditional?

The reason the program terminates without a try/catch is because the exception is actually trying to *propagate upwards/outwards*.

- If the exception can't be caught in the current context/function, then it immediately leaves to the calling context and tries there
  - ▶ And if *that* context can't receive the exception either, then it continues outwards
  - ▶ Since there's nothing above/outside of `main`, that's when it terminates

# Exceptions

## Stack unwinding

That process of leaving the current context, to revert back to the calling context, is called *stack unwinding*.

A bit cliché, but it's not a bug; it's a feature.

If there *is* a catch block in one of the calling contexts, then it can be handled *there*, instead.

- This is both important and valuable, because a local context might not be equipped to fix the problem

This seems like a good place for another example.

# Exceptions

## Catch-all handlers

That last example was missing a single handler.

- Oftentimes, you'll encounter circumstances from which there really is no recovery
  - ▶ Technically, this doesn't *have* to be one of those cases; even after an EOF, you can reset the stream, but that would be contrary to the usage intent
- This presents two special considerations:
  - ▶ We probably can't do anything other than terminate
    - ★ But we should still *exit gracefully*
  - ▶ Whatever we choose to do, something like a complete loss of user input should possibly be handled by the `main`
    - ★ But if we start deferring *all* of our 'special cases' to `main`, we'll eventually need to stop looking for specific cases, particularly if we reach this point because we just need to exit gracefully (irrespective of what the exception is)
- There's a *catch-all* handler: . . .

Let's fix our example, shall we?



## Fancier exceptions

As mentioned earlier, you can throw whatever you need to throw. If you need to throw a monkey, throw a monkey.

But, what if you want better feedback than simply throwing an `int`, or a `double`, or even a `string`, but there's no existing data type that satisfies that need?

- You can define a new class, solely for the purpose of passing along such an exception message, if you wish
- Typically, you'd at least include some form of flag or string message, to hint at what went wrong

Besides the included message/data, simply being a uniquely-chosen type can provide information.

- For example, if you create a `DivisionByZero`, then that's enough to know what happened
  - ▶ Though, if desired, you can still also include member variables for the numerator/denominator

# Fancier exceptions

## Something to be careful about

Note that, just as with any other classes, we *can* include inheritance here as well. e.g.:

- `MathError`
  - ▶ `DivideByZeroError`
    - ★ `ZeroByZeroError` (because we like feeling special)
  - ▶ `NegativeRootError`

If you do this, then be *very* careful of the sequence in which you list the catch blocks.

```
try { /*code*/ }  
catch ( DivideByZeroError ) { /*code*/ }  
catch ( ZeroByZeroError ) { /*code*/ }  
catch ( NegativeRootError ) { /*code*/ }  
catch ( MathError ) { /*code*/ }
```

Because `MathError` is listed last, `DivideByZeroError` and `NegativeRootError` can trigger, but `ZeroByZeroError` *cannot*, as it's listed *after* `DivideByZeroError`!

# Fancier exceptions

## Standard library

It's worth noting the standard library (under the `exception` header) provides for a hierarchy of exceptions.

- The `exception` parent class isn't terribly interesting, but its subtypes are worth looking at (or define your own)
  - ▶ e.g. `runtime_error` accepts a message, to be later retrieved via `.what()`

<http://en.cppreference.com/w/cpp/error/exception>

# Final thoughts on exceptions

Honestly, exceptions can be pretty darn handy, but keep a few things in mind:

- They aren't strictly necessary
  - ▶ For larger, object-oriented systems, you'll probably want a versatile system for error-checking, but for smaller (especially purely procedural) applications (e.g. number-crunching), you might want to skip them
- Be very careful about things like resource management
  - ▶ Closing streams, releasing memory, etc.
  - ▶ There *are* some fancy pointer templates that can help with this
- If you aren't catching the exceptions you're throwing, one might wonder what the point is
- Using exception handling to mask catastrophic errors is like sticking electrical tape over the *check engine* light

## Additional tips

Keep in mind that there are *many* ways to track down persistent bugs.

- There's a Gnu Debugger (GDB), if you're interested in command-line debugging
- Valgrind can be a particularly useful tool for tracking down memory issues
  - ▶ Are you getting inconsistent segfaults? You should check out Valgrind
- Most IDEs allow you to add things like breakpoints, or otherwise aid with debugging

And, of course, the best advice of all for eliminating bugs:

- A well-designed solution with no logic errors to begin with is less work than trying to fix spaghetti

# Questions?

Comments?

- Funky tunes?