

Informed search algorithms

Chapter 3

Outline I

- Informed = use problem-specific knowledge
- Which search strategies?
 - **Best-first search and its variants**
- Heuristic functions?
 - **How to invent them**

Outline II

- Best-first search
- A^*
- Heuristics
- Hill climbing

Kinds of Search problems

- **Type of solution for a given problem**
 - (a) Any solution
 - “How do I get to Toronto? Money/Time is no object!”
 - (b) Optimal solution (best, “good quality”, cheapest,...).
 - “How do I get to Toronto with \$15?”
- **Nature of problem obtained**
 - (a) Finding a path or sequence to solve a problem.
 - path “transforms” start state to goal state
 - E.g., moves needed to solve a Rubik’s Cube?
 - (b) Finding a configuration that is a solution.
 - this single state is everything you need
 - e.g., where to put 8 queens on a board for 8-queen puzzle?

Overview of search algorithms I

- **Basic idea:**

- offline, simulated exploration of state space by generating successors of already-explored states (a.k.a. ~expanding states)
- A strategy is defined by picking the **order of node expansion**

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
```

Previously: tree-search

```
function TREE-SEARCH(problem, fringe) return a solution or failure
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if EMPTY?(fringe) then return failure
    node ← REMOVE-FIRST(fringe)
    if GOAL-TEST[problem] applied to STATE[node] succeeds
      then return SOLUTION(node)
    fringe ← INSERT-ALL(EXPAND(node, problem), fringe)
```

A strategy is defined by picking *the order of node expansion*

Overview of search algorithms II

- **Blind search (uninformed)**
 - Exhaustive search over all configurations
 - “ANY” problem: immediately stop when a solution discovered
 - “Optimal” problem: stop when you are sure best solution found
 - Usually expensive in computation effort!
- **Heuristic Search (informed)**
 - Informed = use problem-specific knowledge
 - “ANY” problems: heuristic helps to find a solution more efficiently
 - heuristics will reduce the number of cases to be looked at
 - although “good” solutions might arise, the best solution is not guaranteed
 - Optimal problems: exhaustive, but can help determine which parts of search tree can be ignored
 - again, heuristic reduces number of cases to investigate
 - a best solution is guaranteed (but computation effort is an issue)

Heuristics

- **Heuristic:** any rule or method that provides some guidance in decision making
 - we use problem-domain specific information in making a decision
 - heuristics vary in the amount of useful information they can lend us
 - e.g., a **stronger heuristic**: don't make any chess move that results in your losing a piece!
 - e.g., a **weaker heuristic** (rule of thumb): knights are best moved into central board positions
- heuristics are often denoted by a functional value: high values denote positive paths, while lower or negative are less promising ones

A heuristic function

- [dictionary] *“A rule of thumb, simplification, or educated guess that reduces or limits the search for solutions in domains that are difficult and poorly understood.”*
 - $h(n)$ = **estimated cost of the cheapest path from node n to goal node.**
 - If n is goal then $h(n)=0$

How to design heuristics? Later..

Heuristics and search

- **Consider a search tree: a given node has a number of children expanded for it (possible all, or just a few)**
 - ideally, we'd like to know which child takes us towards a solution; but this might not be determinable (hence the need for blind search)
- **heuristics permit us to evaluate the children, and select a most promising one**
 - we can even rank them in order of promise
 - this lets us incorporate problem-specific knowledge into the search strategy
 - note: many domains do not admit strong heuristics, so the rating of nodes might be of minimal use (but better than nothing, hopefully!)
- **There are books dedicated to the design of heuristic functions**

Some heuristic search algorithms to consider

- **Optimal (path) search**
 - **Best-first**
 - **A***
- **Non-optimal (local) search**
 - **Hill-climbing**
 - **Beam Search**
 - **Simulated annealing**
 - **Genetic algorithm**

Finding the Best Solution

- aka “optimal search”
 - quality of solution is important
- British museum technique: exhaustively find all paths, then pick the best (e.g., least distance)
 - may use depth-first, breadth-first,... any blind search technique you wish
 - well, search itself isn't important - just exhaustively enumerate all solutions!

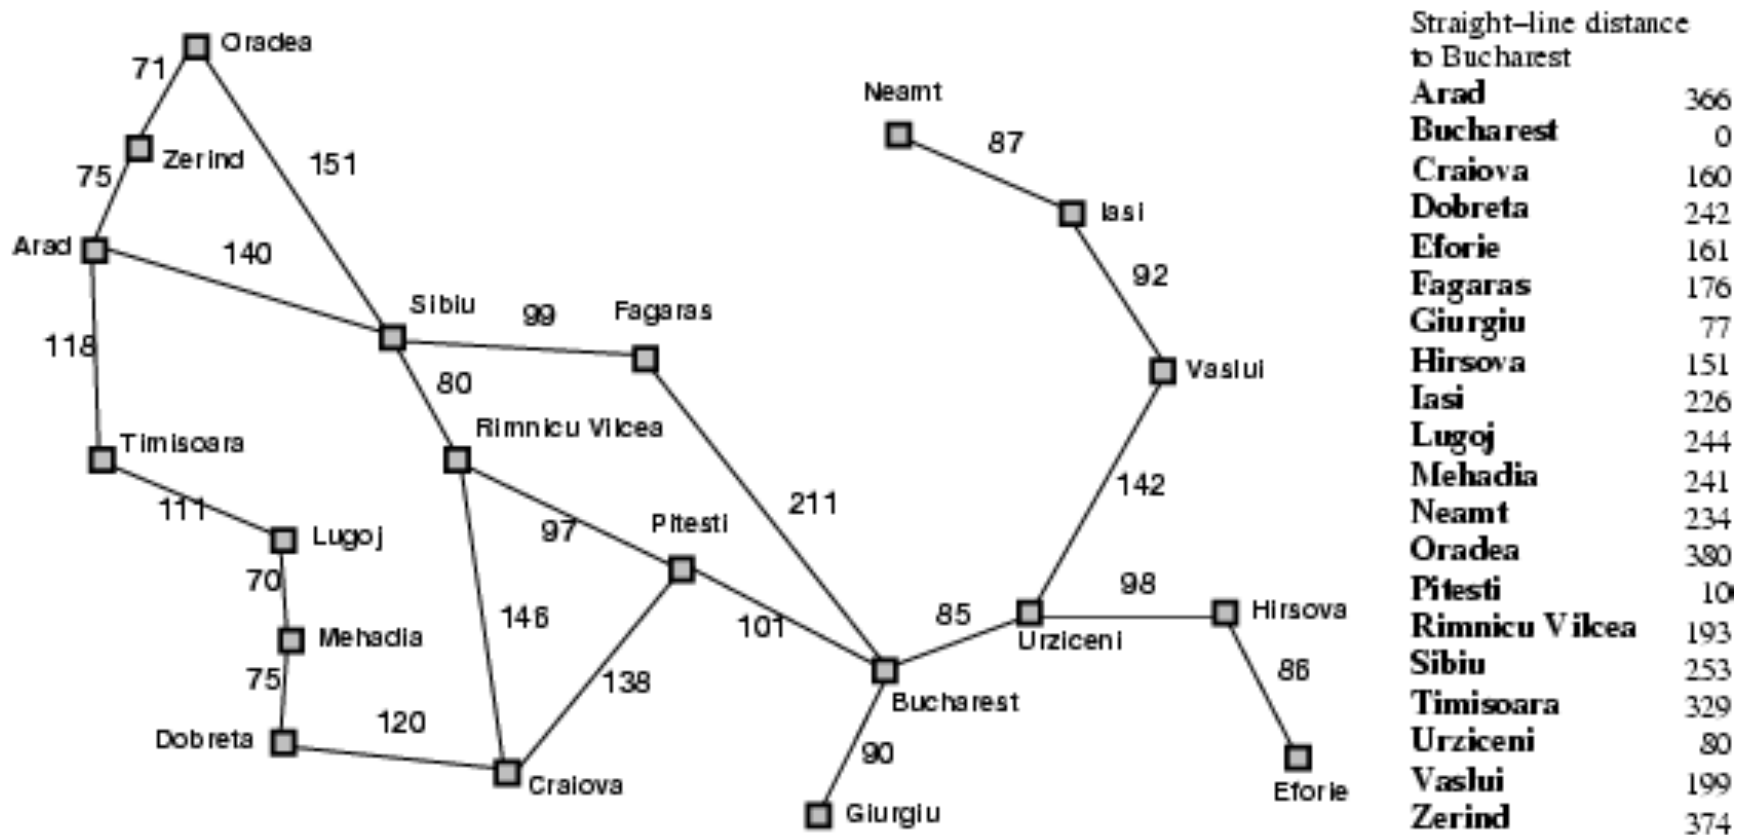
Path problems and Underestimates

- **by adding a guesstimate of the distance remaining for a partial path node, the search can be sped up even more**
 - if you knew exact distance, no search would be necessary
 - if your guess is an overestimate, the problem is that you can no longer use distance information to terminate searches
 - the excess distance value may say a node is too bad to use, when it isn't at all
 - also, any partial path value can't be compared to exact distances with any accuracy
- **how do you make lower-bound estimates?**
 - closer to real values - more accurate the search
 - heuristically

Best-first search

- General approach of informed search:
 - **Best-first search: node is selected for expansion based on an evaluation function $f(n)$**
- Idea: evaluation function measures distance to the goal.
 - **Choose node which *appears* best**
- Implementation:
 - ***fringe* is queue sorted in decreasing order of desirability.**
 - **Special cases: greedy search, A* search**

Romania with step costs in km



Greedy best-first search

- Evaluation function $f(n) = h(n)$ (**h**euristic)
- = estimate of cost from n to *goal*
- e.g., $h_{SLD}(n)$ = straight-line distance from n to Bucharest
- Greedy best-first search expands the node that **appears** to be closest to goal


Greedy best-first search example




Greedy best-first search example



Greedy best-first search example



Greedy best-first search example



Properties of greedy best-first search

- Complete? No – can get stuck in loops, e.g., lasi \rightarrow Neamt \rightarrow lasi \rightarrow Neamt \rightarrow
- Time? $O(b^m)$, but a good heuristic can give dramatic improvement
- Space? $O(b^m)$ -- keeps all nodes in memory
- Optimal? No, same as DF-search

A* search

Idea: avoid expanding paths that are already expensive.

- Evaluation function $f(n) = g(n) + h(n)$

$g(n)$ = cost (so far) to reach n

$h(n)$ = estimated cost to reach the goal from n .

$f(n)$ = estimated total cost of path through n to goal.

A* search


- A* search uses an **admissible** heuristic
 - A heuristic is admissible if it *never overestimates* the cost to reach the goal
 - Are optimistic

Formally:

1. $h(n) \leq h^*(n)$ where $h^*(n)$ is the **true** cost from n
2. $h(n) \geq 0$ so $h(G)=0$ for any goal G .

e.g. $h_{SLD}(n)$ never overestimates the actual road distance

A* search example



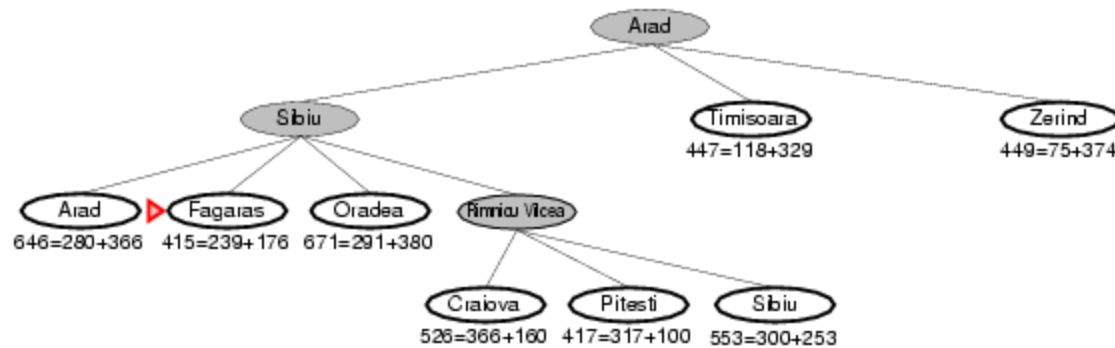
A* search example




A* search example




A* search example



A* search example



A* search example



Optimality of A^*

- Reading assignment

Properties of A*

- Complete? Yes (unless there are infinitely many nodes with $f \leq f(G)$)
- Time? Exponential
- Space? Keeps all nodes in memory
- Optimal? Yes

Improving the memory cost for A*

(Further details found in 1st Edition of your text)

- Algorithm;
 - set cutoff, $h(\text{start node})$
 - i.e., an initial estimate of distance to goal
- Do pure depth-first search, stopping when $f(n) > \text{cutoff}$
 - if succeed, done.
 - if fail, $\text{cutoff} + \text{minimum-amount-by which cutoff-was-exceed}$; iterate step
- This is called “**iterative deepening A***” (IDA)

Iterative Deepening A* (IDA)

- Always finds an optimal solution
- Uses space linear in solution depth
- Is asymptotically no slower than A*
- Assuming a tree structure space, so we don't have to check for cycles
- Or at least, that cycles are few and long
- So, IDA* is about as good as we can do, given an (admissible) heuristic function.

IDA* vs A*

- In practice using a heuristic like the manhattan distance, A* can't solve the 15 puzzle because machine run out of memory, IDA*
- Empirically, IDA* generates more nodes than A*, but surprisingly, it often runs faster!
 - it incurs less overhead per node, and it is easier to implement (since it is essentially depth-first as opposed to breadth-first)
 - e.g., IDA* finds 12 step solution to 8-puzzle problem very quickly expanding 39 nodes
- A lot depends on the problem space
 - For example, in a space where there are only a few values for the heuristic function (e.g., the n-puzzle), IDA* works well, In a case in which each state has a different value (e.g., finding the paths between locations), IDA* will have to expand the square of of the number of nodes A* will expand

Heuristics: defns

- **Admissible heuristic:** the measurement never overestimates the score (distance)
 - otherwise, overestimates may cause good nodes to be skipped, because they are being ignored when they shouldn't be
 - however, if heuristics are too cautious, no useful information is provided. Extreme case: zero for all evaluations (same as using no heuristic!)
- **Consistency:** for distance measurements, a consistent heuristic will give smaller values as you get closer to goal

Admissible heuristics

- A heuristic $h(n)$ is admissible if for every node n , $h(n) \leq h^*(n)$, where $h^*(n)$ is the true cost to reach the goal state from n .
- An admissible heuristic never overestimates the cost to reach the goal, i.e., it is optimistic
- Example: $h_{SLD}(n)$ (never overestimates the actual road distance)
- **Theorem:** If $h(n)$ is admissible, A^* using TREE-SEARCH is optimal

Heuristic functions

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- E.g for the 8-puzzle
 - Avg. solution cost is about 22 steps (branching factor +/- 3)
 - Exhaustive search to depth 22: 3.1×10^{10} states.
 - A good heuristic function can reduce the search process.

Admissible heuristics

E.g., for the 8-puzzle:

- $h_1(n)$ = number of misplaced tiles
- $h_2(n)$ = total Manhattan distance
(i.e., no. of squares from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- $h_1(S) = ?$
- $h_2(S) = ?$

Admissible heuristics

E.g., for the 8-puzzle:

- $h_1(n)$ = number of misplaced tiles
- $h_2(n)$ = total Manhattan distance
(i.e., no. of squares from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- $h_1(S) = 8$
- $h_2(S) = 3+1+2+2+2+3+3+2 = 18$

Inventing admissible heuristics


- Admissible heuristics can be derived from the exact solution cost of a **relaxed** version of the problem:
 - Relaxed 8-puzzle for h_1 : a tile can move **anywhere** then, $h_1(n)$ gives the shortest solution
 - Relaxed 8-puzzle for h_2 : a tile can move **to any adjacent square**.

As a result, $h_2(n)$ gives the shortest solution.

Key point: The optimal solution cost of a relaxed problem is no greater than the optimal solution cost of the real problem.

Inventing admissible heuristics

- Admissible heuristics can also be derived from the solution cost of a subproblem of a given problem.
- This cost is a lower bound on the cost of the real problem.
- Pattern databases store the exact solution to for every possible subproblem instance.
 - The complete heuristic is constructed using the patterns in the DB



Inventing admissible heuristics

- Another way to find an admissible heuristic is through learning from experience:
 - **Experience** = solving lots of 8-puzzles
 - An inductive learning algorithm can be used to predict costs for other states that arise during search.

Learning to search better

- All previous algorithms use *fixed strategies*.
- Agents can learn to improve their search by exploiting the *meta-level state space*.
 - Each meta-level state is a internal (computational) state of a program that is searching in *the object-level state space*.
 - In A^* such a state consists of the current search tree
- A meta-level learning algorithm from experiences at the meta-level.

Local search and optimization

- Previously: systematic exploration of search space.
 - **Path to goal is solution to problem**
- In many optimization problems, the path to the goal is irrelevant; the goal state itself is the solution
- Generally, for some problems path is irrelevant.
 - **E.g 8-queens**
- Different algorithms can be used
 - **Local search**

Example: n -queens

- Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal



Local search and optimization

- Local search= use single current state and move to neighboring states.
- Advantages:
 - Use very little memory
 - Find often reasonable solutions in large or infinite state spaces.
- Are also useful for pure optimization problems.
 - Find best state according to some *objective function*.
 - e.g. survival of the fittest as a metaphor for optimization.

Hill-climbing search

- **It terminates when a peak is reached.**
- Hill climbing does not look ahead of the immediate neighbors of the current state.
- Hill-climbing chooses randomly among the set of best successors, if there is more than one.
- Hill-climbing a.k.a. ***greedy local search***


Hill-climbing search

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                  neighbor, a node

  current ← MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbor ← a highest-valued successor of current
    if VALUE[neighbor] ≤ VALUE[current] then return STATE[current]
    current ← neighbor
```


Hill-climbing search



- Problem: depending on initial state, can get stuck in local maxima.



Hill-climbing example

- 8-queens problem (complete-state formulation).
- Successor function: move a single queen to another square in the same column.
- Heuristic function $h(n)$: the number of pairs of queens that are attacking each other (directly or indirectly).

Drawbacks



- **Ridge** = sequence of local maxima difficult for greedy algorithms to navigate
- **Plateaux** = an area of the state space where the evaluation function is flat.
- Gets stuck 86% of the time.

Hill-climbing variations

- Stochastic hill-climbing
 - Random selection among the uphill moves.
 - The selection probability can vary with the steepness of the uphill move.
- First-choice hill-climbing
 - stochastic hill climbing by generating successors randomly until a better one is found.
- Random-restart hill-climbing
 - Tries to avoid getting stuck in local maxima.

Simulated annealing

Tries to fix the weakness with hill-climbing methods where the search gets stuck in a local maximum.

Basic Idea: Instead of picking the best move, pick a random move; if the successor state obtained by this move is an improvement over the current state, then do it. Otherwise, make the move with some probability < 1 . The probability decreases exponentially with the badness of the move.

Define a temperature function that decreases over time. At each move, compute the current temperature T , and use T to determine the probability with which to allow a move to a worse state. In the limit, T goes to 0 zero at which point the method is doing hill-climbing, hence the probability is proportional to T .

Simulated annealing

- Escape local maxima by allowing “bad” moves.
 - Idea: but gradually decrease their size and frequency.
- Origin; metallurgical annealing
- Bouncing ball analogy:
 - Shaking hard (= high temperature).
 - Shaking less (= lower the temperature).
- If T decreases slowly enough, best state is reached.
- Applied for VLSI layout, airline scheduling, etc.

Simulated annealing

function SIMULATED-ANNEALING(*problem*, *schedule*) **return** a solution state

input: *problem*, a problem
 schedule, a mapping from time to temperature

local variables: *current*, a node.
 next, a node.
 T, a “temperature” controlling the probability of downward steps

current \leftarrow MAKE-NODE(INITIAL-STATE[*problem*])

for *t* \leftarrow 1 **to** ∞ **do**

T \leftarrow *schedule*[*t*]

if *T* = 0 **then return** *current*

next \leftarrow a randomly selected successor of *current*

$\Delta E \leftarrow$ VALUE[*next*] - VALUE[*current*]

if $\Delta E > 0$ **then** *current* \leftarrow *next*

else *current* \leftarrow *next* only with probability $e^{\Delta E / T}$

Local beam search (Reading assignment)

- Keep track of k states instead of one
 - **Initially: k random states**
 - **Next: determine all successors of k states**
 - **If any of successors is goal \rightarrow finished**
 - Else select k best from successors and repeat.
- Major difference with random-restart search
 - **Information is shared among k search threads.**
- Can suffer from lack of diversity.
 - **Stochastic variant: choose k successors at proportionally to state success.**

Search summary

- best search method is problem specific
- heuristics permit us to evaluate the children, and select a most promising one
 - we can even rank them in order of promise
 - this lets us incorporate problem-specific knowledge into the search strategy
 - note: many domains do not admit strong heuristics, so the rating of nodes might be of minimal use (but better than nothing, hopefully!)
- There are books dedicated to the design of heuristic functions